COMPSCI 762 Advanced ML

Github: cwgavin

Sunday, 14 June 2020 10:22 AM

Association Rule Mining

- Frequent itemset generation (support ≥ minsup)
  - Apriori [generate-and-test]
    - supersets of an infrequent itemset must be infrequent (pruned)
  - FP-Tree & FP-Growth [projection-based]
    - scan1: order each row by support, remove infrequent single items
    - scan2: construct FP-tree (compression), build header table
    - for each item in header table
      - construct its conditional pattern base
      - construct its conditional FP-tree, remove infrequent nodes
      - generate frequent patterns
- Rules generation (confidence ≥ minconf)
  - Apriori
    - for same itemset, confidence = / ↓ when #RHS ↑ (pruned)
- Pattern evaluation
  - Interestingness measures
    - confidence, lift / interest, PS, ...
  - Simpson's paradox

Clustering

- Partitional clustering
  - K-means
  - DBSCAN [density-based]
- Hierarchical clustering (dendrogram)
  - Agglomerative
    - Inter-cluster distance
      - min (single link)
      - max (complete link)
      - group average
      - between centroids
      - ward's method (use squared error between point and centroid)
  - Divisive
- Measures of cluster validity
  - entropy (with labels)
  - SSE
  - cohesion - within cluster sum of squares
  - separation - between cluster sum of squares
  - silhouette coefficient

Data Preprocessing

- Data cleaning
  - Missing data
    - imputation
      - global constant
      - average
      - class average
      - most probable value (train a model to predict)
      - matrix decomposition, expectation maximization, ...
  - Noise
    - binning
    - regression
  - Outliers
    - clustering
- Data reduction
  - Dimensionality reduction
    - PCA (only for numeric data)
    - feature selection
      - correlation
        - nominal data: chi-squared test
        - numeric data: pearson's correlation coefficient
      - heuristic search (greedy)
        - Relief (considers all attribtues)
- Data transformation
  - Normalization
    - min-max, z-score, decimal scaling, ...
  - Discretization (continuous → intervals)
    - binning
    - histogram analysis
    - clustering
    - correlation analysis (chi-squared)
    - decision tree
- Imbalanced data
  - undersample majority class
  - oversample minority class
  - cluster-based oversampling
  - SMOTE (create new minority class instance)

Instance-based Learning

- K-Nearest Neighbour (kNN)
  - voronoi diagram
- Support Vector Machine (SVM)

- complexity parameter C (for noise)
- nonlinear
  - Kernel trick
    - calculate distance between points as if transformation is done
- can be applied to regression

Bayesian Learning

- Maximum A Posteriori (MAP)
  - $\arg\max_{h \in H} P(h|D) \approx \arg\max_{h \in H} P(D|h)P(h)$
- Maximum Likelihood (ML)
  - when every hypothesis is equally probable apriori
  - $\arg\max_{h \in H} P(D|h)$
- Minimum description length
  - $\arg\max_{h \in H} P(D|h)P(h)$

    $= \arg\min_{h \in H} -\log_2 P(h) - \log_2 P(D|h)$

    = arg min (length of optimal code for H + ... for D)

    = arg min (complexity of h + #errors made by h)
  - shorter hypothesis is preferred (occam's razor)
  - trade off complexity for #errors -> prevent overfitting
- Bayes optimal classifier (like ensemble)
- Naive bayes classifier
  - laplace smoothing
  - gaussian NBC
  - multinomial NBC
- Bayesian network
  - a directed acyclic graph showing dependencies between attributes

Reinforcement Learning

- Markov decision process (S, A, R, P, γ) - memoryless
  - S: possible states
  - A: possible actions
  - R: reward for state & action
  - P: gives the next state for current state & action
  - γ: discount factor ($0 \leq \gamma < 1$)
- Goal: find optimal action policy $\pi^*: S \rightarrow A$

  $\Leftrightarrow$ find optimal value function V^* (gives value of the state)
- Value Iteration (for learning V^*)
- Q-learning (without knowing P)
  - iteratively simulate markov decision process & update Q matrix

Data Stream

- Sampling
  - reservoir sampling
- Find frequent items
  - lossy counting
  - counting within a sliding window
  - exponential histogram
- Concept drift
  - drift detector
    - CUSUM - alarm when mean differs from 0
    - DDM - monitor error rate
    - ADWIN - based on exponential histogram
    - DDD - ensemble of detectors
- Performance evaluation
  - holdout
  - test-then-train
  - prequential
- Algorithm
  - Hoeffding tree
    - VFDT - with a tie threshold
    - CVFDT - deal with concept drift
  - Hoeffding adaptive tree
    - Hoeffding tree + ADWIN
  - OzaBag
  - Adaptive random forest

Anomaly Detection

- Outliers vs noise
  - Noise: random error or variance, remove first
  - Outliers: violates the normal mechanism that generates the normal data and is interesting
- Statistical (model-based)
  - assume the normal data follow a statistical model (a stochastic model)
  - parametric method
    - univariate data
    - multivariate data
  - non-parametric method
- Proximity-based
  - an object is an outlier if the nearest neighbours are far away
  - distance-based
    - distance/fraction threshold

    - density-based
        - Local Outlier Factor (LOF)
- Clustering-based
    - not belong to any cluster
        - density-based clustering, e.g. DBSCAN
    - far from its closest cluster
        - k-means
    - belong to a small/sparse cluster
        - Cluster-based Local Outlier Factor (CBLOF)
- Classification
    - One-class model, e.g. SVM
    - Semi-supervised
        - clustering-based + one-class model
- Unsupervised
    - Isolation forest
        - Assume anomalies are more easily isolated (shorter path from root)
- Evaluation
    - accuracy is not sufficient
    - use ROC curve (recall vs. FP)